



Ma'at

The Ma'at handbook

How a rumour becomes a verdict: the conversation, the record behind it, the four doors knowledge comes in by, and the rules Ma'at keeps towards the institutions it reads.

Heard something? Check it before you share it.

Contents

1	What Ma'at is	3
2	The conversation, step by step	5
2.1	Opening the widget	5
2.2	Pleasantries, and attempts to steer it	6
2.3	Hearing the rumour	6
2.4	Weighing it against the record	7
2.5	The full path: a rumour that is true	9
2.6	A rumour the record contradicts	11
2.7	The short path: a rumour the record does not settle	13
2.8	After the verdict	15
3	From private answers to public articles	16
4	The record: one pipeline, four doors	18
4.1	The unified ingestion pipeline	18
4.2	The four doors	19
4.3	The scheduler and the country switch	19
4.4	Figures asked for on demand	20
4.5	Proving a source before it is added	20
5	Integrity towards the sites Ma'at reads	21
6	The staff dashboard, page by page	23
7	Limits, honestly stated	30
A	Appendix A · Glossary	32
B	Appendix B · The register in full	33
C	Appendix C · Operations	41
D	Appendix D · Where things live	43

CHAPTER 1

What Ma'at is

Ma'at is a rumour-verification assistant. Somebody has heard something. They bring it to Ma'at. Ma'at weighs it against what reliable institutions have actually published and answers with a verdict and the document behind it, or says honestly that it cannot.

The name is the Egyptian goddess who weighs truth against a feather. The feather is the mark on every page of the product. The promise the product makes to the person who received a rumour, not the one spreading it, is on the landing page in one line:

Heard something? Check it before you share it.

The problem it answers

Rumours travel faster than corrections. A fabricated circular announcing a school holiday, a message that an examinations body will pay everyone's fees, a claim that a ministry has banned something, each reaches a family group before anybody has asked the body concerned. The body has usually spoken; its statement sits on its own website, unread. Ma'at closes that gap: it holds what the institutions have published and lets anyone ask against it in their own words, in their own language, by text or by voice note.

Three principles that govern everything

- 1 Only the record speaks.** Ma'at repeats what a reliable institution has published, and cites it. It never generates a fact of its own. Where the record is silent it says so.
- 2 The gate sits on the judgement, never on similarity.** A passage may be about exactly the right subject and still settle nothing. A verdict is given only when a judgement that a passage supports or contradicts the claim reaches the confidence gate, set at **85%** in Settings.
- 3 Honesty over helpfulness.** "Insufficient evidence" is a real answer, given often, and given with the closest records the visitor can read for themselves.

Where it runs today

Countries

Nigeria and Kenya, both switched on. Coverage is a decision recorded in Settings, not a consequence of which sources happen to exist. A country switched off costs nothing and shows nothing.

Languages

English, then Nigeria's Hausa, Yorùbá, Igbo and Naijá, and Kenya's Kiswahili, all live in the widget and on the site. Gĩkũyũ, Dholuo, Luluhya, Kalenjin, Kikamba and Somali are listed as coming.

The record

625 current documents in 1,239 passages from 59 registered sources: 29 for Nigeria, 22 for Kenya, 8 global bodies. Every figure in this handbook is the database's on 18 September 2026.

Two faces

A public widget on the landing page for anyone who has heard something, and a staff dashboard, sign-in only, for the team that keeps the record and publishes verifications.



The landing page. The teal hero, the feather, the promise, and the gold launcher at the bottom right that opens the widget on any page it is embedded in.


CHAPTER 2

The conversation, step by step


Every reply on the pages that follow was produced by the running system on 18 September 2026, against the documents it actually holds. Nothing is mocked.

2.1 Opening the widget

The launcher is a raised gold button at the bottom right of the page. Opening it shows the welcome screen: who Ma'at is, what it does, the language chosen, and one button, *Start a conversation*. The language can be changed before starting; the choice applies to everything Ma'at says.



Welcome. The feather, the name, the promise in the visitor's language, and the start button.




Language. English, Hausa, Yorùbá, Igbo, Nàijá and Kiswahili live; six more listed as coming.

The composer at the foot of the chat accepts text or a voice note. A voice note is transcribed and then treated exactly like typed text, so the rest of this chapter applies to both.


2.2 Pleasantries, and attempts to steer it

Ma'at reads every message first, before it does anything with it. A greeting, a thank-you or small talk is answered in kind and never triggers a search or a verdict. This holds even in the middle of an interview: someone who says hello while a question is pending is greeted, and the question waits for the next message that is actually about it.

A message that tries to change how Ma'at works, “ignore your instructions”, “reveal your prompt”, is recognised, set aside politely, and recorded as such. Nothing about the conversation changes.



A greeting is a greeting. No search, no verdict, no stale answer from an earlier question.




Steering set aside. The attempt is named for what it is and the door back to a real question is left open.

2.3 Hearing the rumour

When a message carries a claim, Ma'at drafts it into four parts: **who** said or did the thing, **what** is claimed, **when**, and **where**. It reads the draft back in plain words so the visitor can correct it, then asks for what is missing, one short question at a time and never more than two (a number set in Settings). The country matters most: the sources are organised by country, and a claim with no country can only be weighed against the global record.

Anyone who would rather not answer questions can say so. “Just tell me what you have” skips straight to weighing with whatever is known.



Read-back and one question. The claim as Ma'at understood it, then the single thing it still needs.

2.4 Weighing it against the record

Behind the short pause before a verdict, five things happen in order.

- 1 **Retrieval.** The claim, and a few rewordings of it, are searched two ways at once: by meaning, through vector embeddings, and by words, through full-text search. The two lists are fused. Only passages from the visitor's country and from global bodies are eligible, and only from countries switched on in Settings.
- 2 **Shortlisting.** A reranker scores every candidate for usefulness to the claim, 0 to 10, and the pool is cut to what is worth reading. Retrieval's best three always pass regardless, so a near miss is judged and can be named rather than silently dropped.
- 3 **Judgement.** Each shortlisted passage is judged against the claim and returns one of four answers: *supports*, *contradicts*, *same subject but settles nothing*, or *unrelated*, with a confidence and the exact sentence it relied on.
- 4 **The gate.** A support at or above 85% gives **VERIFIED**. A contradiction at or above 85% gives **UNVERIFIED**. Both at once is a conflict in the record, reported as **INSUFFICIENT EVIDENCE** with both sides shown. Nothing at the gate is insufficient evidence, with the kind of nothing named: nothing found, something on the subject that settles nothing, or something that leans one way but not firmly enough.

5

The answer. Written in Ma'at's voice from the decision already made, never the other way round. It names the body, quotes the sentence, gives the date, and links to the original.


Why the gate is on the judgement. On a claim about HIV funding, a passage about the same programme scored very high on similarity and settled nothing. Had the gate been on similarity it would have been called verified. It was called insufficient evidence, which was the truth.

CHAPTER 2 · 2.5


The full path: a rumour that is true

A visitor heard that the President had signed a new National Identity Management Commission Act into law. Here is what Ma'at did with it, end to end.

- 1 **The claim arrived:** "I heard President Tinubu has signed the new NIMC Act 2026 into law. Is it true?" Ma'at read it back and asked when.
- 2 **The visitor answered:** "It happened recently, in Nigeria." The draft was complete: who (the President), what (signed the NIMC Act 2026), when (recent), where (Nigeria).
- 3 **Retrieval** found, among others, NIMC's own press release, read from the Commission's website through the pages door on 17 September.
- 4 **The judge** returned *supports* at 100% and pointed at the sentence.
- 5 **The verdict:** **VERIFIED**, confidence 100.



The verdict card. The verdict, the answer in plain words, and the cited source.



The evidence. The publishing body, the date, the passage with the decisive sentence marked, and the link to the original page.

What Ma'at said

Ma'at

What you heard is supported by the record. The National Identity Management Commission says, “Management Commission (NIMC) proudly announces the signing into law of the National Identity Management Commission (NIMC) Act, 2026 by His Excellency, President Bola Ahmed Tinubu, GCFR.” The News Agency of Nigeria also reports that the NIMC Act 2026 has replaced the 2007 Act and provides a stronger legal framework for identity management. Enough people have raised this for it to stand as a published verdict; the one above is it.

The genuine source, as ingested

National Identity Management Commission press release, read 17 September 2026

PRESIDENT TINUBU SIGNS NIMC ACT 2026 INTO LAW, USHERING IN A NEW ERA OF SECURE DIGITAL IDENTITY FOR NIGERIA. The National Identity Management Commission (NIMC) proudly announces the signing into law of the National Identity Management Commission (NIMC) Act, 2026 by His Excellency, President Bola Ahmed Tinubu.

<https://nimc.gov.ng/press-releases/president-tinubu-signs-nimc-act-2026-into-law...>

The same path, run on a second claim the same day, “Nigeria Customs is recruiting in 2026 and shortlisted people have a physical screening”, also returned **VERIFIED** at 100 against the Customs Service’s own notice: “Attendance at the physical screening and documentation exercise is compulsory, and failure to appear as scheduled may lead to disqualification.”


What makes this genuine. The citation is the publishing body, named exactly as it appears in the register. The quote is the passage the judge relied on, stored word for word, which is why it can open mid-sentence: the passage boundary fell there, and Ma'at does not tidy a source. The link is the body’s own page, not a copy. A reader can open it and check.

CHAPTER 2 • 2.6


A rumour the record contradicts

A fabricated circular said the Federal Ministry of Education had declared a three-week holiday. The Ministry had already answered it. Ma'at found the answer.

- 1 The claim:** "Is it true the Federal Ministry of Education declared a three-week holiday for schools in Nigeria?" The visitor asked to skip the questions.
- 2 The judge** returned *contradicts* at 100%, on the Ministry of Information's statement carrying the Ministry of Education's denial.
- 3 The verdict:** **UNVERIFIED**, confidence 100.



Not supported by the record. The verdict and what the Ministry actually said.



The source. The denial, dated, with the link to the Ministry's page.

Ma'at

This is not supported by the record. The Federal Ministry of Information and National Orientation said the Ministry of Education has “categorically disowned a fabricated circular” that claimed there would be a three-week holiday. They also said all Federal Unity Colleges are expected to resume as planned on 14 September 2026. There is no official announcement of such a holiday. Enough people have raised this for it to stand as a published verdict; the one above is it.

Federal Ministry of Information and National Orientation read 15 September 2026

FME Disowns Fake Circular On Three-week Holliday, Reaffirms Monday Resumption. The Federal Ministry of Education has categorically disowned a fabricated circular currently circulating on social media ... The Minister also urged the public to remain vigilant and to verify information concerning school resumption, Ministry policies, directives and other official announcements through the Federal Ministry of Education's established communication channels before acting on any purported government communication.

<https://fmino.gov.ng/fme-disowns-fake-circular-on-three-week-holiday-reaffirms-monday-resumption/>


The misspelling “Holliday” is the Ministry's own, kept as published. Ma'at never edits a source.

CHAPTER 2 • 2.7

The short path: a rumour the record does not settle

Most rumours are neither confirmed nor denied by anything published. This is the answer given most often, and it is given carefully.


The claim, “WAEC will pay school fees for both vulnerable and non-vulnerable students”, matched nothing from WAEC. The closest record was a News Agency of Nigeria report of a charity paying fees for ten children in Abuja: same subject, different body, and it settles nothing. Ma'at says so, shows what it found, and asks before looking further.



Asking first. What was looked for, that it was not enough, and the offer to ask trusted online sources, taken up only on a yes.

What “yes” does

On consent, Ma'at asks the statistical APIs configured for the country, the World Bank, the WHO Global Health Observatory and UNHCR, for figures that speak to the claim, stores whatever comes back as documents under that country, and weighs the claim again with the shelf now larger. If



Insufficient evidence. The honest verdict, with the closest records labelled as such, each with its body, date, passage and link, so the visitor can judge.

nothing bears on it, it says so in one line. Every lookup and every consent, given or refused, is kept on record.

Why the near miss is shown

A reply of “nothing relevant” when a closely related record exists would be false. The judge’s category *same subject, settles nothing* exists for this, and the retrieval’s best three passages always reach the judge so that category can be reached. The visitor sees up to three closest records, labelled “Closest record” rather than “Cited source”, ranked by how well they matched, each with its body, date, passage and link, and decides for themselves whether one was what they meant.

CHAPTER 2 · 2.8

After the verdict

A copy of the document

Where the cited document was uploaded by staff and marked *shared*, Ma'at offers a copy: "I have the document itself. Would you like a copy?" On a yes the file arrives in the conversation with its title, issuer, name and size, and a download icon. A refusal is honoured and not asked again for that question. Documents read from feeds, APIs or pages are never offered as copies; they are linked to the publisher's page instead.

The same rumour, raised again

Every claim is matched against rumours already held. The same rumour in different words is folded into one, and counts as one more mention. If the rumour has been published, the visitor is told so alongside the verdict. If it was last raised more than **30 days** ago, Ma'at asks once whether this is something new before folding it in, because the same words can describe a different event a year later.

Memory, and the lack of it

- A conversation lasts as long as the widget is on the page. Reloading starts a new one. The screen and the server never disagree about what has been said.
- A conversation left untouched for **30 minutes** drops its half-finished interview. The next message starts clean.
- The visitor's raw words are kept for **30 days** (set in Settings) and then expire; the paraphrase, verdict and citations are kept.
- A visitor cookie counts distinct browsers across conversations for publication, and nothing else. It holds no identity.
- Rate limits: 20 questions per visitor per hour, 2,000 a day across the platform, both set in Settings.

CHAPTER 3


From private answers to public articles

A rumour is answered privately to each visitor until enough different people have raised it. Then its status changes, and Ma'at starts telling people so.

Publication turns on the **reporter count**: how many different browsers raised the rumour, counted from the visitor cookie. Nine messages from one determined person are one reporter; three different browsers within minutes are three. At **3 reporters** (set in Settings) the rumour's status becomes *Published* automatically, with its current verdict and citations, unless staff have marked it withheld.

- From then on, a visitor who raises the same rumour is told that Ma'at has published on it before, alongside the verdict.
- A withheld rumour stays private however many raise it; the count keeps rising so the decision can be revisited.
- The three worked examples in Chapter 2 crossed the threshold while this handbook was being made, each raised from more than three browsers, and were published automatically. The Rumours page below shows them.

What “published” means today. The status, the referral in the widget, and the count on the dashboard are live. The public article page itself, and the list of verifications on the landing page, are not yet built; the landing page shows sample articles. This is listed in Chapter 7.




Rumours as staff see them. Verdicts by day and by share for the chosen range, then the rumours most recently raised: statement, verdict and confidence, the body the verdict rests on, mentions, reporters, status and when it was last seen.

CHAPTER 4

The record: one pipeline, four doors

Everything Ma'at can cite came in through one of four doors and went down one pipeline. The doors differ only in how a document is fetched. From the moment it is in hand, every document is treated identically.



The flow of knowledge. Four doors in, one pipeline, one record, one retrieval. Figures asked for on demand join the record as documents like any other.

4.1 The unified ingestion pipeline

- 1 Parse.** PDF (layout mode), DOCX, ODT, RTF and plain text for uploads; the article text for feeds, APIs and pages. A file with no readable text, a scanned PDF for instance, is refused with a plain message rather than stored mute.
- 2 Clean.** Invisible characters some publishers' editors sprinkle through prose are removed; text that was double-encoded on the publisher's side ("â€™" for an apostrophe) is put back the way it was typed. Both were found in real documents and both broke search.
- 3 Split.** Sections first, so a heading stays attached to its passage, then sentences. A passage that reads fine and cites nothing is the failure this avoids.
- 4 Enrich.** Each passage gets one line situating it in its document before embedding, so a fragment reading "the rate is 7.5%" is findable by the levy it belongs to. Capped at 200 passages a document.
- 5 Embed.** Vectors from the model provider. Identical bytes reuse vectors already stored, so a re-poll costs nothing.
- 6 Store.** Every document is versioned. A re-upload or a corrected article becomes a new version; the old one is retired, never erased. Answers come from the newest.

- 7 File.** A national source's documents go under its country. A global feed's story is read for which covered country it names; a story naming two is filed global.

The safety property. Parsing and embedding both happen before the previous version is retired, and both raise rather than degrade. Done the other way round, the old version would be retired, replaced with nothing, and the upload would report success while the document went mute.

4.2 The four doors, in order of preference

ORDER	DOOR	WHAT IT IS	TODAY
1	Public API	A structured endpoint the body exposes. Includes the WordPress REST API most public bodies' sites carry, which returns complete articles with dates and links. Read by a schema on the source: dot paths to items and fields, fallbacks, nulls, dates normalised, and a raw sample kept on the first success so the schema is written against what the API really returns.	42 sources
2	Feed	RSS or Atom the body maintains. Fetched with conditional requests, so an unchanged feed costs one small reply. Entries without a stable identifier are dropped rather than duplicated.	10 sources
3	Public web pages	The body's own news pages, for official bodies with neither of the above. Read under the rules in Chapter 5. Through a headless browser only where the site draws its news with JavaScript.	7 sources
4	Direct upload	Files staff add on the Documents page, attributed to a body on the register, shelved under a country, and marked shared or private.	Ma'at itself

The door is decided by asking the site what it offers, in that order (4.5), never by assumption. If a body later adds a feed or an API, it moves up a door.

4.3 The scheduler and the country switch

A poller runs as a daemon beside the API. Every five minutes it looks for sources whose cadence has elapsed (six hours for national bodies, daily for statistical platforms) and reads each one with the reader for its door. Every poll is recorded, including the ones that found nothing, because "when did this last work" is the question a quiet source raises. A source that fails five times in a row stops being polled until somebody looks at it; one source failing never stops the others.

Settings holds the one switch that matters: which countries are on. A country switched off is **not polled, not searched, not asked and not shown**. Its documents stay on record and are counted on the Settings page; nothing else spends a request or a token on it. Switching it on starts all four at once. Adding a country switched off is how it is prepared: its sources are registered and proven before Ma'at answers for it.

4.4 Figures asked for on demand

Statistical warehouses publish numbers, not announcements; there is nothing to poll. The World Bank, the WHO Global Health Observatory and UNHCR are registered as sources marked *asked on demand*. When a claim needs a figure and the visitor consents, the relevant series are asked for the visitor's country, the answer is written as a short quotable document under that country, dated the day it was asked, and kept for a day before the same question would go back out. The second person asking is answered from the shelf.

4.5 Proving a source before it is added

`discover_source <address>` asks a site, in order: is there a WordPress REST API; is there a feed announced on the page or at the usual paths; are there sitemaps; does the news page yield a link that reads as a complete, dated article. Each answer is verified by actually reading it. The command recommends a door and prints the configuration to store. After seeding, the first poll must bring documents in before the source is considered proven. Every row in Appendix B passed both.

On 15 September 2026, 72 official sites in Nigeria and Kenya were probed this way. 38 were added. The rest are listed in Chapter 7 with the reason, so the same ground is not walked twice.

CHAPTER 5

Integrity towards the sites Ma'at reads

Ma'at reads other institutions' websites. How it does so is a matter of principle, enforced in code, not left to intention. Every rule below is in `knowledge/pages.py` and has a test.

Who may be read

The pages door is for official and public-interest bodies only: ministries, agencies, regulators, commissions, public broadcasters, the state news agency, state-owned bodies whose statements the public depends on, and recognised humanitarian bodies such as the national Red Cross. Private newspapers, television, radio, blogs and aggregators are never read this way; their business is the page. Who qualifies is a register decision made by staff, and the line is held there.

How it is read

- 1 robots.txt first.** Fetched before anything else on the site and obeyed to the letter, including `Crawl-delay`. A robots file that cannot be read because the server errored *closes* the door for that poll; nothing is fetched. A missing robots file is taken, by the web's convention, to mean public pages are fine.
- 2 Identified.** Every request, from every door, carries the user agent `MaatBot/1.0 (+https://maatverify.com; rumour verification)`. A body that would rather not be read can see who is asking and say so.
- 3 Slowly.** One request at a time, at least two seconds apart per site, longer if robots asks. At most 20 new articles a poll, so a backlog is read over days rather than in one night. At most 300 links considered from one listing.
- 4 Whole articles, once.** The article is separated from navigation, sidebar and footer. A page that does not read as a complete article, a menu, a gallery, a stub, is not kept. A page already held is never fetched again.
- 5 Attributed.** Stored with the body's name, the story's own title, its date as the page states it, and its address. Cited that way, with a link to the original.
- 6 Nothing from behind a door.** No login, no paywall, no cookies, no forms, no images, media or fonts, no comments, no following links off the site.
- 7 A browser only where the page needs one.** Some official sites draw their news with JavaScript. For those, and only those, a source is marked `render` and its pages are fetched through headless Chromium that still announces itself as Ma'at, keeps the pacing, and loads no images or media. Same extractor, same rules.

What Ma'at will not do

Pretend to be a person

Two sites answer every request with a redirect only a human's browser with cookies resolves. That is a gate. Ma'at does not put on a browser's user agent to pass it. Those sites are out until they open a door.

Read over a broken connection

Twelve Kenyan government sites serve https with certificates that do not validate. A rumour-checking service cannot cite a page it cannot prove came from that body. They are out until they fix it.

Push past a refusal

A "too many requests" answer on the first request is taken as the answer. No retries, no rotating addresses.

Guess a date, edit a source

Dates come from the page's own markup or not at all; an invented date on a citation is worse than none. Text is stored as published, misspellings included.

We read what a public body has put in public, the way a careful citizen reads a notice board: openly, slowly, without touching anything, and leaving when asked.

What this costs is stated plainly in Chapter 7: the bodies that cannot be read today, and why.


CHAPTER 6

The staff dashboard, page by page

Behind sign-in, the team sees the same record the widget answers from, and every figure on it is counted from the database, not typed in.

Sign-in


Staff sign in through a dialog over the site, or are sent to it when they open a dashboard address unsigned. The dashboard code only downloads once past the guard. Sign-out returns to the landing page.



Sign-in. A dialog in the site's style; the same account signs into Django's admin.

Dashboard


The front page for the chosen range (7, 30 or 90 days): rumours **weighed**, the share **cited** with a verified or unverified verdict, **documents in**, each with an arrow against the period before. The digest card sums this period's verdicts and how many were published. Activity shows rumours raised by day. Top sources ranks the bodies holding most of what Ma'at can cite, with their own marks. The channels band counts the messages that reached the widget in the period; the channels not yet connected are named as coming, with no figure invented for them.




The front page on 18 September 2026. Five rumours weighed that week, three of them cited, 625 documents in, 35 messages through the widget, and Kenya Revenue Authority holding the largest share of the record.

Sources

The register: every body Ma'at cites, as cards (20 a page) or a compact list (50 a page). Each shows the body's mark, its country or "Global", its door, its address, cadence, last check, document count, verification status and health. A country lens sits inside the search box and offers only the countries switched on in Settings. The panel at the top states the four doors and their rules. *Add a source* opens a three-step form: the body, the door, the connection.



Cards. Each body with its own mark, door badge, address and health.



List. The whole register at a glance, on the same frosted surface.

Ma'at

staffcheck

Staff

Dashboard

Rumours

Conversations

Sources

Documents

Settings

Log out

Add a source

← Back to sources

1 The body

What a citation will name, and the mark shown beside it.

Name

The publishing body as it should appear in a citation. Write it as they write it.

World Health Organization

What it publishes

One line. It appears on the card, so the register reads as more than a list of addresses.

Press releases and situation reports

Scope

A body whose word applies anywhere is global. Otherwise name the country.

Global

Logo

An https address for the body's own mark. Left empty, a monogram stands in.

https://example.org/logo.svg

2 How it connects

Pick one. The server prefers an API, then a feed, then pages, and checks the site for the better door first.

Direct upload

Staff add the files by hand. For bodies that publish no feed and run no API.

RSS or Atom feed

A feed the body publishes, polled on a schedule. New entries are fetched and read.

Public API

A documented endpoint the body offers openly. Structured, and the most reliable door.

Public web pages

An official body's own news pages, read where it offers neither API nor feed. Whole articles, by its robots rules, slowly.


Adding a source. Name, what it publishes, scope, then the door, then the connection the door needs.

Documents


The upload folder at the top takes PDF, DOCX, ODT, RTF and TXT, attributes each file to a body from the register, shelves it under a country or global, and can mark it shared on the way in. Below, the shelves: one per country plus global, closed by default with their counts in the heading, remembered when opened, opened by a search that matches inside them. Each row shows the file type, the title, the publisher's mark and name, the date, passages, size and the shared/private toggle. Twenty documents a shelf, paged.

Heard something? Check it before you share it.

26




Documents. The upload folder and the shelves, closed until wanted.



A shelf open. Publisher marks beside each document, the shared toggle at the end of the row.

Settings


The single source of truth for what Ma'at covers and how it decides. **Countries:** the list, each with its source and document counts and its switch, new ones added off. **Verdict thresholds:** the confidence gate (85), reporters before publication (3), follow-up questions (2). **Rate limits and retention:** questions per visitor per hour (20), the daily cap (2,000), raw-text retention (30 days).




Settings. Countries with their switches, then the thresholds the engine reads on every turn.

Rumours, Conversations, Profile

Rumours shows verdicts by day and by share for the chosen range, then the rumours most recently raised with the body each verdict rests on (Chapter 3). Conversations is a placeholder today (Chapter 7). Profile shows the signed-in member and offers sign-out.



Conversations. Present, not yet built out.



Profile. The member, the role, sign-out.

CHAPTER 7

Limits, honestly stated

What Ma'at cannot do today, and why. Each is either a site's decision, a fact about the web, or work not yet done.

Bodies probed and not added

BODY	WHY NOT
Ministry of Health (Kenya), KNBS, Kenya News Agency, IEBC, National Police Service, KNEC, Interior, Foreign Affairs, EPRA, NTSA, KWS, NCIC	Their https certificates do not validate. Ma'at does not read over a connection it cannot verify. Revisit when they fix it.
NiMet, Pharmacy and Poisons Board (Kenya)	Every request answered with a redirect only a browser with cookies resolves: a bot challenge. We do not pretend to be a browser.
Nigeria Police Force	Rate-limits every request (HTTP 429).
FIRS, Federal Ministry of Education, FCCPC, KeNHA	Down or timing out when probed. Retry later.
Central Bank of Nigeria, FRSC, NECO, NYSC, NIHSA, NSIB, NDLEA, NNPC, NMDPRA, JAMB	No API, no feed, and no news listing that yields readable articles, in a browser or not: brochure and shell pages only.
INEC	Press pages carry 150 characters of text; the release itself is an embedded image or file.
Kenya Power	Its newsroom links to an articles listing that does not answer, as a page or in a browser.
MyGov (Kenya)	Renders in a browser, but none of its links reads as a complete article; it is a weekly paper laid out as pages.
National Bureau of Statistics	Brochure site; no news listing found.
Our World in Data, UNICEF Data Warehouse	Registered from the original specification with base addresses and no adapter yet; kept on the register switched off.

Not yet built

- **Public article pages.** A published rumour has a status, a referral in the widget and a count on the dashboard, but no public page yet; the verifications listed on the landing page are sample content.
- **Conversations page.** A placeholder; the conversations exist in the database.

- **Saving a source from the form.** The form validates and previews; sources are added to the register file and seeded. Auth blocks and schemas are edited there too.
- **Channels.** Only the widget exists and is counted; WhatsApp, X and Facebook are named on the dashboard band as coming, with no figures.
- **Median response time** on the front page, replaced by documents in until there is timing data.
- **A “yes, that is what I meant” step** on the closest records shown for an unsettled rumour.
- **Native-speaker review** of the translations.

Facts about the corpus

- Many feeds carry summaries, not whole articles. A summary is still the body's own words, dated and attributable, which is what a citation needs.
- A national broadcaster's world news is filed under its country, because that is who reported it to that country.
- The record is young: 625 documents on 18 September 2026. Insufficient evidence will be the common answer until it grows, and that is the correct answer.

APPENDIX A

Glossary

TERM	MEANING
Claim	One visitor's statement of a rumour, drafted into who, what, when, where.
Rumour	The canonical form many claims are folded into. Carries the verdict, the confidence, the mention and reporter counts, and the status: collecting, published, withheld.
Mention	One raising of a rumour by one claim.
Reporter	One distinct browser that raised the rumour, counted from the visitor cookie. What publication turns on.
Verdict	Verified (green), Unverified (amber), Insufficient evidence (grey). Never red.
Confidence gate	The judgement confidence a support or contradiction must reach before it becomes a verdict. 85 by default.
Evidence	A passage judged against a rumour, with the judgement, the score and the quoted sentence.
Source	A body on the register: name, door, address, country or global, cadence, logo, and for APIs an auth block and a schema.
Door	How a source's publications arrive: public API, feed, public web pages, direct upload.
Document	One publication, versioned, filed under a country or global, marked shared or private.
Passage (chunk)	A piece of a document with its heading path, its enrichment line and its vector. What retrieval returns and the judge reads.
Coverage	A country on the Settings list, with its switch.
Live lookup	Asking the on-demand APIs for a figure, on consent, and storing the answer.
Published	A rumour raised by enough different reporters. Its status, and what the widget tells the next person who raises it. The public page for it is not yet built.

APPENDIX B

The register in full

Every source as it stands on 18 September 2026. Each national row was read live before it was kept.

Nigeria

BODY	DOOR	ADDRESS	DOCS	STATE
Federal Ministry of Agriculture and Food Security Food security, farm inputs and agricultural policy	API	agriculture.gov.ng/wp-json/wp/v2/posts	11	polled
Federal Ministry of Finance Fiscal statements, budgets and clarifications	API	www.finance.gov.ng/wp-json/wp/v2/posts	12	polled
Federal Ministry of Health and Social Welfare Ministerial statements, outbreak responses and health policy	API	health.gov.ng/wp-json/wp/v2/posts	10	polled
Federal Ministry of Information and National Orientation Official statements and clarifications from the federal government	API	fmino.gov.ng/wp-json/wp/v2/posts	20	polled
Federal Radio Corporation of Nigeria News from the public broadcaster	Feed	radionigeria.gov.ng/feed/	10	polled
Humanitarian Data Exchange Datasets filed under Nigeria by WFP, WHO, UNHCR, IOM and NGOs	API	data.humdata.org/api/3/action/package_search?fq=groups:nga	5	polled
National Agency for Food and Drug Administration and Control Product recalls, fake and substandard medicine alerts, public notices	API	nafdac.gov.ng/wp-json/wp/v2/posts	12	polled
National Emergency Management Agency Disaster alerts, flood warnings, relief operations and situation updates	API	nema.gov.ng/wp-json/wp/v2/posts	10	polled

National Health Insurance Authority Health insurance notices and clarifications	API	www.nhia.gov.ng/wp-json/wp/v2/posts	10	polled
National Identity Management Commission NIN enrolment, identity notices and press releases	Pages · browser	nimc.gov.ng/press-releases/	10	polled
National Pension Commission Pension notices and regulations	API	www.pencom.gov.ng/wp-json/wp/v2/posts	11	polled
National Primary Health Care Development Agency Immunisation campaigns, vaccine notices and primary care	API	nphcda.gov.ng/wp-json/wp/v2/posts	12	polled
News Agency of Nigeria Reports from the state news agency	API	nannews.ng/wp-json/wp/v2/posts	40	polled
Nigeria Centre for Disease Control and Prevention Outbreak alerts, public health advisories and situation reports	Pages	ncdc.gov.ng/news-recent	21	polled
Nigeria Customs Service Import and export notices, seizures and recruitment announcements	Pages · browser	customs.gov.ng/news	5	polled
Nigeria Immigration Service Passport, visa and border notices	API	immigration.gov.ng/wp-json/wp/v2/posts	10	polled

Nigerian Civil Aviation Authority Aviation safety directives and press releases	Pages	ncaa.gov.ng/media/press-releases/	2	polled
Nigerian Communications Commission SIM registration, NIN linkage and network notices	Feed	www.ncc.gov.ng/rss.xml	10	polled
Nigerian Electricity Regulatory Commission Tariff orders and electricity notices	API	nerc.gov.ng/wp-json/wp/v2/posts	5	polled
Nigerian Television Authority News reports from the public broadcaster	Pages	www.nta.ng/news	40	polled
Nigeria Security and Civil Defence Corps Security advisories, arrests and infrastructure protection	API	nscdc.gov.ng/wp-json/wp/v2/posts	12	polled
Our World in Data Curated indicators for Nigeria, each citing its own source	API	ourworldindata.org/grapher/?csvType=filtered&country=NGA	0	off
Securities and Exchange Commission Nigeria Investor alerts and warnings about unregistered schemes	Feed	www.sec.gov.ng/feeds/news.rss	6	polled

State House, Nigeria Statements and releases from the Presidency	API	statehouse.gov.ng/wp-json/wp/v2/posts	16	polled
UNHCR Refugees and displaced people hosted in Nigeria, and those who left it. Asked on demand...	API · on demand	api.unhcr.org/population/v1/population/?coa=NGA	0	on demand
UNICEF Data Warehouse Child health, education, nutrition and water for Nigeria	API	sdmx.data.unicef.org/ws/public/sdmxapi/rest/data/UNICEF,GL	0	off
West African Examinations Council, Nigeria Examination timetables, results releases and notices	Feed	www.waecnigeria.org/rss.xml	10	polled
WHO Global Health Observatory Health statistics for Nigeria: mortality, immunisation, disease burden. Asked on demand...	API · on demand	ghoapi.azureedge.net/api/?\$filter=SpatialDim eq 'NGA'	0	on demand
World Bank Indicators for Nigeria: economy, population, health. Asked on demand when a claim needs a...	API · on demand	api.worldbank.org/v2/country/NGA/indicator/	1	on demand

Kenya

BODY	DOOR	ADDRESS	DOCS	STATE
Central Bank of Kenya Monetary policy, currency notices and bank licensing	API	www.centralbank.go.ke/wp-json/wp/v2/posts	10	polled
Communications Authority of Kenya SIM registration, network and broadcasting notices	Feed	www.ca.go.ke/rss.xml	9	polled
Directorate of Criminal Investigations Investigations, arrests and public warnings	Feed	www.dci.go.ke/rss.xml	2	polled
Humanitarian Data Exchange Datasets filed under Kenya by WFP, WHO, UNHCR, IOM and NGOs	API	data.humdata.org/api/3/action/package_search?fq=groups:ken	0	polled
Kenya Airports Authority Airport operations and travel advisories	API	www.kaa.go.ke/wp-json/wp/v2/posts	10	polled
Kenya Broadcasting Corporation News from the public broadcaster	API	www.kbc.co.ke/wp-json/wp/v2/posts	31	polled
Kenya Bureau of Standards Product standards, recalls and counterfeit alerts	API	www.kebs.org/wp-json/wp/v2/posts	9	polled
Kenya Medical Research Institute Research findings and public health statements	API	www.kemri.go.ke/wp-json/wp/v2/posts	13	polled
Kenya Meteorological Department Weather warnings, seasonal outlooks and climate statements	Pages	meteo.go.ke/news/	4	polled
Kenya Red Cross Society Emergency response, appeals and situation updates	API	redcross.or.ke/wp-json/wp/v2/posts	10	polled
Kenya Revenue Authority Tax notices, enforcement and clarifications	Pages	www.kra.go.ke/news-center/press-release	79	polled

Ministry of Agriculture and Livestock Development Food security, fertiliser subsidies and livestock notices	API	kilimo.go.ke/wp-json/wp/v2/posts	10	polled
Ministry of Education, Kenya School calendars, examinations and education notices	Feed	www.education.go.ke/rss.xml	10	polled
National Drought Management Authority Drought early-warning bulletins and relief	API	ndma.go.ke/wp-json/wp/v2/posts	10	polled
Our World in Data Curated indicators for Kenya, each citing its own source	API	ourworldindata.org/grapher/?csvType=filtered&country=KEN	0	off
Social Health Authority Registration, contributions, benefits and clarifications	API	www.sha.go.ke/wp-json/wp/v2/posts	12	polled
The National Treasury, Kenya Budget statements, public debt and fiscal notices	Feed	www.treasury.go.ke/rss.xml	10	polled
The Presidency, Kenya Statements and releases from State House	API	www.president.go.ke/wp-json/wp/v2/posts	10	polled

UNHCR Refugees and displaced people hosted in Kenya, and those who left it. Asked on demand...	API · on demand	api.unhcr.org/population/v1/population/?coa=KEN	0	on demand
UNICEF Data Warehouse Child health, education, nutrition and water for Kenya	API	sdmx.data.unicef.org/ws/public/sdmxapi/rest/data/UNICEF,GL	0	off
WHO Global Health Observatory Health statistics for Kenya: mortality, immunisation, disease burden. Asked on demand...	API · on demand	ghoapi.azureedge.net/api/?\$filter=SpatialDim eq 'KEN'	0	on demand
World Bank Indicators for Kenya: economy, population, health. Asked on demand when a claim needs a...	API · on demand	api.worldbank.org/v2/country/KEN/indicator/	0	on demand

Global

BODY	DOOR	ADDRESS	DOCS	STATE
Humanitarian Data Exchange Datasets from WFP, WHO, UNHCR, IOM and NGOs	API	data.humdata.org/api/3/action/	0	off
Our World in Data Curated cross-country indicators with their sources cited, as CSV per chart	API	ourworldindata.org/grapher/	0	off
UNHCR Refugees, internally displaced people and asylum, by country of asylum or origin. Asked...	API · on demand	api.unhcr.org/population/v1/	0	on demand
UNICEF Data Warehouse Child health, education, nutrition and water. SDMX, readable as CSV	API	sdmx.data.unicef.org/ws/public/sdmx-api/rest/	0	off
UN News All stories: humanitarian, political, health	Feed	news.un.org/feed/subscribe/en/news/all/rss.xml	53	polled
WHO Global Health Observatory Health statistics by country: mortality, immunisation, disease burden. Asked on demand...	API · on demand	ghoapi.azureedge.net/api/	0	on demand
WHO Regional Office for Africa Press releases	Feed	afro.who.int/rss/featured-news.xml	22	polled
World Bank Indicators API: economy, population, health. No authentication. Asked on demand when a...	API · on demand	api.worldbank.org/v2/	0	on demand

APPENDIX C

Operations

COMMAND	WHAT IT DOES
<code>python scripts/dev_up.py</code>	Starts the backend (gunicorn, https, port 8765), the frontend (Vite, port 5174) and the poller daemon, inside the backend virtualenv. <code>dev_down.py</code> stops them.
<code>manage.py poll_feeds --loop</code>	The poller daemon: ticks every five minutes, polls what is due. <code>--force</code> polls everything once; <code>--every N</code> changes the tick.
<code>manage.py discover_source URL</code>	Probes a site for an API, a feed, or readable pages; <code>--render</code> probes through the browser; <code>--file</code> takes a list.
<code>manage.py seed_sources</code>	Creates or updates the register from <code>knowledge/fixtures/register.json</code> , schemas and auth included.
<code>python scripts/fetch_logos.py</code>	Fetches each body's official mark from Wikipedia or its own site and normalises it to the register's 192px canvas.
<code>python -m playwright install chromium</code>	Installs the headless browser used for sources marked <code>render</code> .
<code>manage.py test</code>	The backend suite: 243 tests on 18 September 2026, run with the model provider offline so every stage takes its deterministic path.

Environment

`backend/.env` holds the database URL, the secret key, the allowed hosts and origins, the front-end URL, and the OpenAI key and model names for the everyday model (reading the visitor, the interview), the careful model (judging evidence, writing the answer) and the reranker. `AI_OFFLINE=1` runs every stage without a provider. Uploaded and archived documents live under `backend/media/`, outside the repository.

Timings and thresholds that are not in Settings

CONSTANT	VALUE	WHERE
Poller tick	5 minutes plus jitter	<code>poll_feeds</code>
Failures before a source stops being polled	5	<code>feeds.py</code>
Pages: pause between requests to one site	2 seconds, or robots' crawl-delay	<code>pages.py</code>
Pages: new articles per poll / links per listing	20 / 300	<code>pages.py</code>
Pages: what counts as an article	500 characters with a date, or 1,200 without	<code>pages.py</code>

Browser: page arrival / settle wait	25 s / 8 s	pages.py
API sample kept on first success	3 items, 1 nested item, 200-character strings, 16 KB	connectors.py
On-demand figures kept before re-asking	24 hours	lookup.py
Reranker floor / passages always judged	3.0 / 3	engine.py
Same rumour, vector distance	0.10	engine.py
Recurrence question after	30 days	engine.py
Conversation goes cold after	30 minutes	engine.py
Upload size cap	25 MB	knowledge/api.py

APPENDIX D

Where things live

PATH	WHAT IT HOLDS
backend/core	Soft-delete base model, ISO catalogues (countries, currencies, time zones), current-user context.
backend/accounts	Staff profile, session API, styled sign-in and sign-out pages.
backend/ai	The model provider and every model-calling stage: interpreter, interview, reranker, judge, answer, enrichment, embeddings. Each with an offline path.
backend/knowledge	The record: models, parsing, services (the pipeline), feeds, connectors (schema-driven API), pages (the pages door and the browser), lookup (on-demand APIs), geography (country filing), discovery, the register fixture, the poller command.
backend/verification	The engine (the conversation), retrieval, the chat API, the visitor cookie, models for conversations, claims, rumours, mentions, evidence, articles, and the dashboard figures.
backend/appsettings	Thresholds and rate limits, country coverage and its switch, and the one rule every other module asks: <code>switched_on()</code> .
frontend/src/widget	The embeddable widget: launcher, welcome, language, chat, bubbles with links and attachments, verdict card, API client.
frontend/src/landing	The public site: hero, articles, footer, brand.
frontend/src/auth	The sign-in dialog, the guard, the session.
frontend/src/dashboard	The staff dashboard: sections for overview, rumours, sources, source form, documents, settings, profile; shared components including the publisher mark and the pager.
frontend/src/i18n	Messages and the six live locales.
docs/	<code>verification-engine.md</code> (the agreed engine design), <code>sources.md</code> (the doors policy and register), this handbook.



When there is no source, Ma'at says so.

Ma'at handbook · produced from the running system on 18 September 2026